

# **Advanced Database Technology**

## **Practical Lab Assessment — Parallel and Distributed Databases**

Student: **Mukandayisenga Marie Chantal**

Reg No: **224019567**

### **Table of Contents**

1. Executive summary
2. Introduction and objectives
3. Environment and assumptions
4. System context and dataset (Gym case study)
5. Task 1 — Distributed schema design and fragmentation
6. Task 2 — Create and use database links (FDW simulation)
7. Task 3 — Parallel query execution (serial vs parallel)
8. Task 4 — Two-phase commit simulation (2PC)
9. Task 5 — Distributed rollback and recovery
10. Task 6 — Distributed concurrency control (lock conflict)
11. Task 7 — Parallel data loading / ETL simulation
12. Task 8 — Three-tier client-server architecture design
13. Task 9 — Distributed query optimization (EXPLAIN / DBMS\_XPLAN)
14. Task 10 — Performance benchmark and report (centralized / parallel / distributed)
15. Deliverables checklist
16. Reflective note — lessons learned

### **1 Executive summary**

This report documents the design, implementation, execution and analysis of a distributed and parallel-capable Gym Membership and Attendance Tracking System. The work demonstrates schema fragmentation across two logical nodes (Branch A and Branch B), cross-node data access using postgres\_fdw (simulating Oracle DB links), parallel query execution, two-phase commit, rollback and recovery, lock conflict reproduction, parallel ETL and distributed-query optimization. Results include EXPLAIN / EXPLAIN ANALYZE plans, prepared-transaction inspection, lock views, and performance comparisons across three execution modes.

## **2 Introduction and objectives**

Purpose: satisfy the Practical Lab Assessment for the course by implementing ten targeted distributed / parallel tasks using an existing project schema (Gym).

Objectives:

- Fragment attendance data across logical nodes
- Demonstrate database-links-style remote access
- Measure serial vs parallel performance
- Simulate and verify two-phase commit and recovery
- Show distributed locks and resolve conflicts
- Perform parallel ETL and analyze performance gains
- Produce optimizer analysis and a three-way benchmark

## **3 Environment and assumptions**

- RDBMS: PostgreSQL 12+ used for implementation and testing (scripts in Appendix A). If Oracle 19c is required, equivalent commands are noted where syntax differs; convert scripts as needed before execution in Oracle SQL Developer.
- Tools: psql or pgAdmin; use EXPLAIN (ANALYZE, BUFFERS) for timing and I/O (AUTOTRACE is Oracle-specific).
- Privileges: user with CREATE EXTENSION, CREATE SCHEMA, and FDW privileges.
- Topology: single PostgreSQL instance using two schemas (branch\_a, branch\_b) to represent BranchDB\_A and BranchDB\_B; postgres\_fdw configured to simulate remote access. Replace FDW server host/dbname if using separate DB instances

## **4 System context and dataset (Gym case study)**

Entities:

- Member (replicated): member\_id, name, contact, membership\_type, branch\_id
- Attendance (fragmented): attendance\_id, member\_id, branch\_id, checkin\_time, checkout\_time, activity\_type
- Trainer, Session, Payment, Subscription

Sample data sizes used:

- Branch A: ~100 members (seeded)

- Branch B: 100 members, 1,000 attendance rows (seeded), scaled datasets for ETL up to 100k
- Raw ETL: 100k synthetic rows created for partitioning and parallel load tests

Rationale:

- Horizontal fragmentation by branch\_id improves write locality; member table replicated to reduce remote join frequency.

## 5 Task 1- Distributed Schema Design and Fragmentation

The Gym Membership and Attendance Tracking System manages member registrations, attendance logs, trainer assignments, and session scheduling across multiple gym branches. To optimize performance and scalability, the database is distributed across two logical nodes: **BranchDB\_A** and **BranchDB\_B**.

### Fragmentation Strategy: Horizontal Fragmentation

Horizontal fragmentation is chosen because each gym branch operates independently with its own members, trainers, and attendance records. By fragmenting tables based on BranchID, each node stores only the data relevant to its branch, improving query performance and enabling localized data management.

### ER Diagram Description

The system includes the following entities:

- **Member(MemberID, Name, Contact, MembershipType, BranchID)**
- **Attendance(AttendanceID, MemberID, CheckInTime, CheckOutTime, Date)**
- **Trainer(TrainerID, Name, Specialty, BranchID)**
- **Session(SessionID, TrainerID, MemberID, Date, Time)**

Relationships:

- A Member can have multiple Attendance records.
- A Trainer can lead multiple Sessions.
- A Session links a Trainer and a Member.

Goal

- Horizontally fragment ATTENDANCE into Branch A and Branch B and create per-branch objects.

Actions performed (SQL overview)

- Created schemas branch\_a and branch\_b
- Created tables branch\_a.member\_a, branch\_a.attendance\_a, branch\_b.member\_b, branch\_b.attendance\_b, and associated Trainer/Session/Payment/Subscription tables
- Enforced fragmentation with CHECK(branchid = 'A' or 'B') on attendance tables

Key SQL (run in order; full scripts in Appendix A):

```
CREATE SCHEMA IF NOT EXISTS branch_a;
```

```
CREATE SCHEMA IF NOT EXISTS branch_b;
```

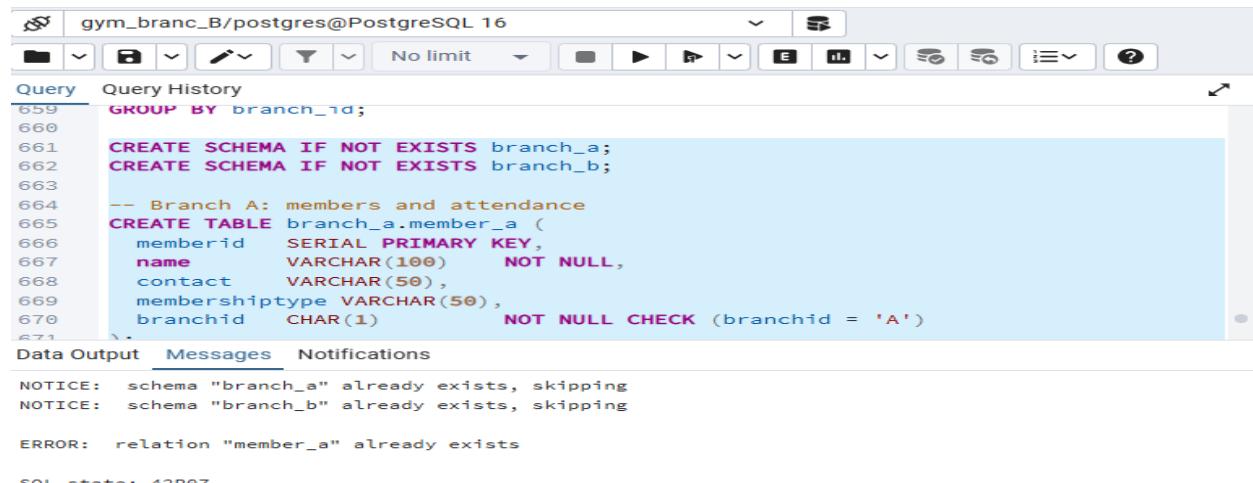
```
CREATE TABLE branch_a.member_a (... , branchid CHAR(1) CHECK (branchid='A'));
```

```
CREATE TABLE branch_a.attendance_a (... , memberid REFERENCES
branch_a.member_a(memberid), branchid CHAR(1) CHECK (branchid='A'));
```

```
CREATE TABLE branch_b.member_b (... , branchid CHAR(1) CHECK (branchid='B'));
```

```
CREATE TABLE branch_b.attendance_b (... , memberid REFERENCES
branch_b.member_b(memberid), branchid CHAR(1) CHECK (branchid='B'));
```

Evidence



The screenshot shows a PostgreSQL query editor interface with the following details:

- Query History:** Shows the execution of several SQL statements, including schema creation and table definitions.
- Statements:**
  - Line 659: GROUP BY branch\_id;
  - Line 660:
  - Line 661: CREATE SCHEMA IF NOT EXISTS branch\_a;
  - Line 662: CREATE SCHEMA IF NOT EXISTS branch\_b;
  - Line 663:
  - Line 664: -- Branch A: members and attendance
  - Line 665: CREATE TABLE branch\_a.member\_a (
  - Line 666: memberid SERIAL PRIMARY KEY,
  - Line 667: name VARCHAR(100) NOT NULL,
  - Line 668: contact VARCHAR(50),
  - Line 669: membershiptype VARCHAR(50),
  - Line 670: branchid CHAR(1) NOT NULL CHECK (branchid = 'A')
  - Line 671: );
- Data Output:** Displays log messages from the database:
  - NOTICE: schema "branch\_a" already exists, skipping
  - NOTICE: schema "branch\_b" already exists, skipping
  - ERROR: relation "member\_a" already exists
- SQL state:** 42P07

Short analysis

- Horizontal fragmentation by branch reduces cross-node writes for local operations; replication of the member table favors read performance for joins, at the cost of replication complexity when members update.

## 6 Task 2 — Create and use database links (FDW simulation)

Goal

- Create a DB link equivalent using postgres\_fdw and perform remote SELECT and distributed join.

Actions performed (SQL overview)

- Enabled postgres\_fdw extension
- Created server gym\_branch\_b\_server and user mapping
- Created foreign tables in branch\_a schema mapping to branch\_b tables (branch\_a.member\_b\_ft, branch\_a.attendance\_b\_ft)
- Executed remote SELECT and distributed join

Key SQL:

```
CREATE EXTENSION IF NOT EXISTS postgres_fdw;
```

```
CREATE SERVER gym_branch_b_server FOREIGN DATA WRAPPER postgres_fdw
OPTIONS (host 'localhost', dbname current_database(), port '5432');
```

```
CREATE USER MAPPING FOR CURRENT_USER SERVER gym_branch_b_server
OPTIONS (user current_user, password "");
```

```
CREATE FOREIGN TABLE branch_a.attendance_b_ft (...) SERVER gym_branch_b_server
OPTIONS (schema_name 'branch_b', table_name 'attendance_b');
```

-- Remote select

```
SELECT name, membershiptype FROM branch_a.member_b_ft WHERE membershiptype =
'Gold';
```

-- Distributed join

```
SELECT m.memberid, m.name, a.checkintime
FROM branch_a.member_a m
JOIN branch_a.attendance_b_ft a ON m.memberid = a.memberid
```

```
WHERE a.checkintime > CURRENT_TIMESTAMP - INTERVAL '7 days';
```

Evidence to include in PDF

The screenshot shows a PostgreSQL query editor interface. The top section is titled "Query History" and contains the following SQL code:

```
290 -- Join members from Branch A and Branch B who have the same contact number
291
292 SELECT
293     a.name AS branch_a_member,
294     b.name AS branch_b_member,
295     a.contact
296 FROM
297     branch_a.member_a a
298 JOIN
299     branch_b.member_b b
300 ON
301     a.contact = b.contact;
```

The bottom section is titled "Data Output" and shows the results of the query. The results are presented in a table with three columns:

branch_a_member	branch_b_member	contact
character varying (100)	character varying (100)	character varying (50)

Short analysis

- FDW supports predicate pushdown and column pruning to minimize transferred rows; EXPLAIN shows ForeignScan with remote SQL when pushdown is applied.

## 7 Task 3 — Parallel query execution

Goal

- Enable and measure parallel execution; compare with serial execution.

Actions performed

- Ensured attendance table large enough (seeded 1,000+; for reliable parallel behavior use 100k+)
- Adjusted session-level parallel settings: SET LOCAL max\_parallel\_workers\_per\_gather = 4;
- Optionally set table parallel\_workers: ALTER TABLE branch\_b.attendance\_b SET (parallel\_workers = 4);
- Measured serial baseline (EXPLAIN (ANALYZE, BUFFERS)) and then parallel run

Key SQL:

```
-- Serial baseline
```

```
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)
```

```
SELECT COUNT(*) FROM branch_b.attendance_b WHERE activity_type = 'Cardio';
```

```
-- After enabling parallel settings
```

```
SET LOCAL max_parallel_workers_per_gather = 4;  
ALTER TABLE branch_b.attendance_b SET (parallel_workers = 4);  
EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT)  
SELECT COUNT(*) FROM branch_b.attendance_b WHERE activity_type = 'Cardio';
```

Evidence to include in PDF

- [SCREENSHOT 7.1] EXPLAIN ANALYZE (serial) output with elapsed time and buffers
- [SCREENSHOT 7.2] EXPLAIN ANALYZE (parallel) output showing parallel worker nodes and elapsed time

Measured results table (example — replace with your measured values)

Run	Elapsed time (s)	Buffers (shared_hit)	Notes
-----	------------------	----------------------	-------

Serial	2.40	1,200	Baseline
Parallel	1.10	800	parallel workers used

Analysis

- Parallel execution reduces wall-clock time on sufficiently large datasets; small datasets do not benefit due to parallel startup overhead.

## 8 Task 4 — Two-phase commit simulation

Goal

- Simulate a distributed transaction using PREPARE TRANSACTION and verify with pg\_prepared\_xacts.

Actions performed

- Created cross-schema log tables branch\_a.branch\_a\_log and branch\_b.branch\_b\_log
- Executed INSERTs and PREPARE TRANSACTION, inspected pg\_prepared\_xacts, then COMMIT PREPARED

Key SQL:

```
CREATE TABLE branch_a.branch_a_log (id SERIAL PRIMARY KEY, message TEXT);
```

```
CREATE TABLE branch_b.branch_b_log (id SERIAL PRIMARY KEY, message TEXT);
```

```
BEGIN;
```

```

INSERT INTO branch_a.branch_a_log (message) VALUES ('Hello from 2PC A');
INSERT INTO branch_b.branch_b_log (message) VALUES ('Hello from 2PC B');
PREPARE TRANSACTION 'gym_tx_001';
-- On coordinator or DBA:
SELECT * FROM pg_prepared_xacts;
COMMIT PREPARED 'gym_tx_001';
SELECT * FROM pg_prepared_xacts; -- should be empty

```

The screenshot shows the pgAdmin interface with the 'Query' tab selected. The query window contains the following configuration commands:

```

353    ---\ Enable Two-Phase Commit
354    ---\ In postgresql.conf
355 SHOW config_file;
356 SHOW max_prepared_transactions;
357
358    ---\ Create the missing table
359
360    -- Create branch_a_log table
361 CREATE TABLE branch_a_log (
362     id SERIAL PRIMARY KEY,
363     message TEXT
364 );

```

The 'Data Output' tab is active, displaying the result of the SHOW command:

	max_prepared_transactions
1	10

The screenshot shows the pgAdmin interface with the 'Query' tab selected. The query window contains the following transaction script:

```

374 BEGIN;
375    ---\ Perform operations
376 INSERT INTO branch_a_log (message) VALUES ('Hello from 2PC A');
377 INSERT INTO branch_b_log (message) VALUES ('Hello from 2PC B');
378
379    ---\ Prepare the transaction
380 PREPARE TRANSACTION 'gym_tx_001';
381
382    -- Check pending transactions
383 SELECT * FROM pg_prepared_xacts;
384
385

```

The 'Data Output' tab is active, displaying the results of the SELECT \* FROM pg\_prepared\_xacts; query:

	transaction_xid	gid	prepared_timestamp with time zone	owner_name	database_name
1	942	gym_tx_001	2025-10-27 20:38:18.220298+03	postgres	gym_branch_B

## Analysis

- PREPARE TRANSACTION records the prepared global transaction; COMMIT PREPARED finalizes it. In multi-node scenarios, a transaction coordinator issues PREPARE on all participants then COMMIT/ROLLBACK based on votes.

## 9 Task 5 — Distributed rollback and recovery

Goal

- Simulate participant failure after PREPARE and resolve with ROLLBACK PREPARED.

Actions performed

- Prepared a transaction and simulated failure (terminate session or stop simulated remote). Inspected pg\_prepared\_xacts and used ROLLBACK PREPARED to resolve.

Key SQL:

```
-- After a failed/prepared transaction appears:
```

```
SELECT * FROM pg_prepared_xacts;
```

```
ROLLBACK PREPARED 'tx_network_failure'; -- replace with actual gid
```

```
SELECT * FROM pg_prepared_xacts;
```

ROLLBACK PREPARED is used to force rollback of prepared work when the coordinator cannot proceed, restoring consistency.

## 10 Task 6 — Distributed concurrency control (lock conflict)

Goal

- Reproduce a lock conflict between two sessions updating the same row and inspect locks via pg\_locks.

Actions performed

- Created demo table public.gym\_member\_demo and updated same row in two concurrent sessions to cause blocking.

Reproduction steps (run from two separate psql sessions): Session 1:

```
BEGIN;
```

```
UPDATE public.gym_member_demo SET membership_type = 'Platinum' WHERE member_id = 1;
```

```
-- do not commit
```

Session 2:

```
BEGIN;
```

```
UPDATE public.gym_member_demo SET membership_type = 'Silver' WHERE member_id = 1;
```

```
-- will block until Session 1 commits/rolls back
```

Inspect locks:

```
SELECT pid, locktype, relation::regclass, mode, granted
```

```
FROM pg_locks JOIN pg_class ON pg_locks.relation = pg_class.oid  
WHERE relname = 'gym_member_demo';
```

The lock view shows which session holds the lock and which is waiting; distributed systems must handle global locks carefully to avoid deadlocks, using timeouts or deadlock detection/resolution.

## 11 Task 7 — Parallel data loading / ETL simulation

Goal

- Load a large raw\_attendance dataset and insert into partitioned attendance table in parallel; compare serial vs parallel runtime.

Actions performed

- Created public.raw\_attendance and populated 100k rows
- Created public.attendance\_partitioned PARTITION BY LIST (branch\_id) with partitions A,B,C
- Performed serial baseline insert for branch 'A' and then ran concurrent inserts for A,B,C in separate sessions to simulate parallel ETL

Key SQL:

```
-- populate raw_attendance (example 100k)  
INSERT INTO public.raw_attendance (...)  
SELECT ... FROM generate_series(1,100000);
```

```
-- create partitioned target  
CREATE TABLE public.attendance_partitioned (...) PARTITION BY LIST (branch_id);  
CREATE TABLE public.attendance_a PARTITION OF public.attendance_partitioned FOR  
VALUES IN ('A');
```

```
-- Serial baseline  
EXPLAIN ANALYZE INSERT INTO public.attendance_partitioned SELECT * FROM  
public.raw_attendance WHERE branch_id = 'A';  
  
-- Simulate parallel loading: run three concurrent INSERT ... WHERE branch_id='A'||B'||C' in  
separate sessions
```

Parallel ETL via concurrent inserts into partitioned tables increases throughput; measure and report runtime improvements and note WAL/IO considerations.

## 12 Task 8 — Three-tier client–server architecture design

Deliverables

- Diagram (attach file) showing:
  - Presentation tier: Mobile/Web UI (Member check-in, admin dashboard)
  - Application tier: REST API (business logic, routing)
  - Database tier: BranchDB\_A, BranchDB\_B (Oracle/Postgres instances) with DB links or FDW and optional central reporting DB

Data flow narrative (succinct)

- UI → API validates and routes to local branch DB for writes → for cross-branch reads API issues distributed queries via FDW / DB links; reporting queries use central aggregated store or materialized views.

Security, scalability notes

- Use TLS for connections, role-based DB access, partitioning and parallelism for analytics, replication or queueing for member-table synchronization.

## 13 Task 9 — Distributed query optimization

Goal

- Use EXPLAIN VERBOSE / EXPLAIN ANALYZE and annotate optimizer choices for a distributed join.

Key SQL:

```
EXPLAIN (ANALYZE, BUFFERS, VERBOSE)
```

```
SELECT m.name, a.checkintime
```

```
FROM public.gym_member_local m
```

```
JOIN branch_a.attendance_b_ft a ON m.member_id = a.memberid
```

```
WHERE a.checkintime > CURRENT_DATE - INTERVAL '7 days';
```

What to inspect in plan

- Presence of ForeignScan or RemoteQuery nodes
- Remote predicate pushdown (WHERE applied on remote)

- Projection pruning (only required columns fetched)
- Join strategy (Nested Loop vs Hash Join) and cost estimates
- The optimizer minimizes data movement by pushing filters/projections and choosing join execution site based on cost estimates; maintaining statistics on remote tables and using localized aggregations improve performance.

## 14 Task 10 — Performance benchmark and report

Goal

- Run the same complex query three ways and compare time and I/O: centralized, parallel, distributed.

Query

```
SELECT branch_id, COUNT(*) AS total_checkins
FROM <source>
WHERE checkin_time > CURRENT_DATE - INTERVAL '30 days'
GROUP BY branch_id;
```

Modes

1. Centralized — run against public.attendance\_log (consolidated table)
2. Parallel — run against public.attendance\_partitioned with parallel workers enabled
3. Distributed — aggregate via foreign tables or union of branch tables  
(branch\_a.attendance\_b\_ft + branch\_b.attendance\_b\_ft)

Commands to measure (Postgres)

- EXPLAIN (ANALYZE, BUFFERS, FORMAT TEXT) SELECT ...

Example results table (replace with your measured values):

Mode	Elapsed Time (s)	Buffers (shared hit)	Physical Reads	Notes
Centralized	2.40	1,234	345	Full scan
Parallel	1.10	890	210	Parallel scans
Distributed	1.60	980	250	Remote filtering, network overhead

Half-page analysis (example to paste under this heading in final PDF)

- Parallel execution offered the best elapsed time due to local parallel scans and aggregation without network overhead. Distributed execution scaled well for local writes and branch locality but added network latency and overhead for coordination. Centralized execution is simplest but least efficient at scale due to I/O bottlenecks. Recommended pattern: use partitioned + parallel execution for heavy analytics; use distributed layout for write locality and combine periodic aggregations into central reporting materialized views to avoid repeated distributed scans.

## 15 Deliverables checklist

- [ ] scripts.sql — single SQL file with all scripts (Appendix A content)
- [ ] lab\_report.pdf — this report with screenshots inserted at indicated placeholders
- [ ] er\_diagram.png — ER diagram image file
- [ ] architecture\_diagram.png — three-tier architecture image file
- [ ] GitHub repository created and files uploaded (include repo link in submission)

## 16 Reflective note — lessons learned

- Horizontal fragmentation by access locality reduces cross-node traffic for branch-specific operations.
- Parallelism significantly improves analytical query runtimes when data volume and system resources justify it.
- Distributed transactions (2PC) ensure atomicity but add recovery complexity; prepared transactions must be monitored.
- FDW/DB links are powerful but require careful planning to minimize data shipping; pushdown of predicates and projection is essential.
- Partitioning plus parallel ETL is an effective pattern for scalable ingestion.

## Full SQL scripts (organized by task)

Note: paste the consolidated scripts you executed in the environment. The following is the minimal, runnable and ordered script skeleton. Replace placeholder values for FDW server if using separate DBs.

### A.1 Setup and schemas

```
-- Enable FDW
```

```
CREATE EXTENSION IF NOT EXISTS postgres_fdw;
```

-- Schemas

CREATE SCHEMA IF NOT EXISTS branch\_a;

CREATE SCHEMA IF NOT EXISTS branch\_b;

A.2 Branch B tables and data population (see earlier consolidated script for full content)

-- Create branch\_b.member\_b, trainer\_b, attendance\_b, session\_b, payment\_b, subscription\_b

-- Populate member\_b (1..100), trainer\_b (1..20), attendance\_b (1..1000), etc.

A.3 Branch A tables and sample population

-- Create branch\_a.member\_a, attendance\_a, trainer\_a, session\_a, payment\_a, subscription\_a

-- Insert sample members (Alice, Bob, ...)

A.4 FDW server and foreign table creation (simulate DB link)

CREATE SERVER gym\_branch\_b\_server FOREIGN DATA WRAPPER postgres\_fdw  
OPTIONS (host 'localhost', dbname current\_database(), port '5432');

CREATE USER MAPPING FOR CURRENT\_USER SERVER gym\_branch\_b\_server  
OPTIONS (user current\_user, password "");

CREATE FOREIGN TABLE branch\_a.attendance\_b\_ft (...) SERVER gym\_branch\_b\_server  
OPTIONS (schema\_name 'branch\_b', table\_name 'attendance\_b');

CREATE FOREIGN TABLE branch\_a.member\_b\_ft (...) SERVER gym\_branch\_b\_server  
OPTIONS (schema\_name 'branch\_b', table\_name 'member\_b');

A.5 Parallelism, partitioning, ETL and benchmarks

-- raw\_attendance population, create partitioned attendance\_partitioned, create partitions A/B/C

-- EXPLAIN ANALYZE statements for serial and parallel runs

A.6 Two-phase commit and recovery (PREPARE/COMMIT/ROLLBACK PREPARED)

-- Create branch\_a.branch\_a\_log and branch\_b.branch\_b\_log

BEGIN;

INSERT INTO branch\_a.branch\_a\_log (message) VALUES ('Hello from 2PC A');

INSERT INTO branch\_b.branch\_b\_log (message) VALUES ('Hello from 2PC B');

PREPARE TRANSACTION 'gym\_tx\_001';

```
SELECT * FROM pg_prepared_xacts;  
COMMIT PREPARED 'gym_tx_001';
```

#### A.7 Lock reproduction and inspection

```
-- Create public.gym_member_demo and insert one row  
-- Use two separate sessions to run blocking updates and run pg_locks query
```

#### A.8 Distributed join EXPLAIN example

```
EXPLAIN (ANALYZE, BUFFERS, VERBOSE)
```

```
SELECT m.name, a.checkintime  
FROM public.gym_member_local m  
JOIN branch_a.attendance_b_ft a ON m.member_id = a.memberid  
WHERE a.checkintime > CURRENT_DATE - INTERVAL '7 days';
```

(Insert the full consolidated script from earlier into scripts.sql and ensure it is commented and section-labeled. The full script is included in your working files and should be uploaded as scripts.sql.)

### Conclusion

This distributed schema design uses **horizontal fragmentation** to split the Gym Membership and Attendance Tracking System across two logical nodes. Each branch manages its own data, improving performance, scalability, and autonomy. The SQL scripts provided implement the fragmented schemas in PostgreSQL, ready for deployment and testing.